

CricMetrics Pro: Complete Machine Learning Topics & Methodology Guide

Simple & Clear Technical Guide | Case Study no. 102 Reference

About this Document: This guide explains all the Machine Learning concepts, algorithms, math formulas, and testing methods used in the **CricMetrics Pro** sports analytics project. It is written in simple, clear, and direct English so that every student, examiner, and data scientist can easily understand **WHAT** each technique does, **WHEN** it is used, **WHY** it was chosen, and **HOW** it works.

1. Project Overview & Machine Learning Tasks

In modern T20 cricket, franchises (like IPL teams) spend ₹100+ Crores in player auctions. Choosing players based on emotions, raw reputation, or simple batting averages often leads to expensive mistakes.

CricMetrics Pro solves this by analyzing **260,920 real deliveries from 1,095 IPL matches (2008–2024)** covering **619 qualified professional cricketers**. We set up 4 core Machine Learning tasks:

ML Task	What It Predicts / Outputs	Business Goal	Models Used	Best Model & Result
1. Talent Classification	Talent Tier Code (0, 1, 2)	Group players into Developing, Core, or Elite tiers	• K-Nearest Neighbors (KNN) • Random Forest Classifier • Logistic Regression • Decision Tree	K-Nearest Neighbors (KNN) (95.16% Accuracy) & Random Forest (94.35%)
2. Unsupervised Clustering	Tactical Style (5 Archetypes)	Group athletes by how they actually play, not just "Batter/Bowler"	• K-Means Clustering (k=5 clusters)	5 Clear Tactical Archetypes (Anchor, Finisher, Fast Bowler, Spinner, All-Rounder)
3. Dimensionality Reduction	2D Coordinates (X, Y)	Compress 29 stats into an interactive 2D map	• Principal Component Analysis (PCA)	Top 2 Components explain 60.79% of all variation

2. Data Preprocessing & Feature Engineering

Before training models, raw cricket data must be cleaned, transformed, and rescaled.

2.1 One-Hot Encoding (primary_role)

- WHAT is it?**
Converting text labels (like "Specialist Batter" or "All-Rounder") into 0 and 1 columns that computers can do math with.
- WHEN is it used?**
Applied during initial data preparation in `src/cricket_data_pipeline.py` before feeding data to any model.

- **WHY is it necessary?**

Computers only calculate numbers. If we gave roles numbers like `1 = Batter, 2 = All-Rounder, 3 = Bowler`, the model would mistakenly think a Bowler is "greater than" a Batter. One-hot encoding creates separate True/False (1/0) switches for each role. We drop the first column (`drop_first=True`) to prevent mathematical redundancy (the "dummy variable trap").

- **HOW is it done?**

```
`python
```

```
df_encoded = pd.get_dummies(df, columns=['primary_role'], drop_first=True, dtype=float)
```

```
`
```

Creates 4 binary columns: `primary_role_Bowling Specialist`, `primary_role_Specialist Batter`, `primary_role_Specialist Bowler`, and `primary_role_Squad Batter`. If all 4 are 0, the player is an `All-Rounder`.

2.2 Standard Scaling (Z-Score Normalization)

- **WHAT is it?**

Rescaling every continuous column so that its average becomes 0 and its standard deviation becomes 1:

$$z = \frac{x - \mu}{\sigma}$$

- **WHEN is it used?**

Applied strictly **after** the 80/20 train/test split. The scaler learns only from training data (`X_train`) to avoid cheating (data leakage), and then scales both training and test data.

- **WHY is it necessary?**

In cricket, `total_runs` can be 5,000+ while `economy_rate` is only 7.5. Without scaling, distance calculations in KNN, K-Means, and PCA would be 99% dominated by runs, completely ignoring bowling stats. Scaling puts every skill on equal ground.

- **HOW is it done?**

```
`python
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
`
```

Saved as `models/scaler.joblib`.

2.3 Cricket Composite Indices (Domain Feature Engineering)

Raw career counting totals (like total runs) heavily favor older players who played 200 matches over younger stars who played 30 matches. To fix this, we created 3 rate-based impact indices:

1. Batting Impact Index:

Combines Batting Average, Strike Rate, Boundary %, and Death-Overs Strike Rate (overs 16–20).

2. Bowling Impact Index:

Combines Economy Rate, Bowling Strike Rate (balls per wicket), and Dot Ball %.

3. Clutch Match-Winner Index:

Rewards players who step up under pressure: Player of the Match awards, match-winning fifties, and 3+ wicket hauls.

3. Supervised Learning: Talent Tier Classification Models

Classification models place players into **3 talent tiers**:

- **Tier 1 (Elite / Marquee)**: Overall Rating ≥ 80.0 benchmark (Star match-winners).
 - **Tier 2 (Core / Star)**: Overall Rating 68.0 to 79.9 benchmark (Reliable tournament starters).
 - **Tier 0 (Developing / Squad)**: Overall Rating < 68.0 benchmark (Emerging players & squad backups).
-

3.1 K-Nearest Neighbors (KNN, $k=7$) — 🏆 TOP ACCURACY

- **WHAT is it?**

A simple, intuitive classifier: to classify a player, it looks at the **7 most similar players** in cricket history and picks the most common tier among them.

- **WHEN is it used?**

Used in Tab 2 and Tab 6 for talent tier classification.

- **WHY was it chosen?**

1. **Highest Accuracy**: Achieved **95.16% test accuracy** (top among all classifiers).

2. **Matches Real Cricket Scouting**: Scouts naturally evaluate new players by comparing them to similar past players ("He bowls and bats just like Hardik Pandya").

- **HOW does it work & results?**

- Measures Euclidean distance across all 29 scaled stats:

$$d(p, q) = \sqrt{\sum (p_i - q_i)^2}$$

- Takes a distance-weighted vote of the 7 closest neighbors ($k=7$).

- **Test Accuracy**: 95.16% | **Macro F1-Score**: 0.9501 | **Precision**: 0.9601 | **Recall**: 0.9421 .

3.2 Random Forest Classifier — 🏆 PRODUCTION TIER ENGINE

- **WHAT is it?**

An ensemble of **150 decision trees** voting on which tier a player belongs to.

- **WHEN is it used?**

Powers the Tier Classification engine and the **Feature Importance Leaderboard** in Tab 6.

- **WHY was it chosen?**

1. **Very High Accuracy**: 94.35% on test data, and highest cross-validation score (94.75%).

2. **Feature Importance (Explainability)**: It tells franchise owners which stats matter most using Gini impurity.

- **HOW does it work & results?**

- Implemented using `RandomForestClassifier(n_estimators=150, max_depth=10, random_state=42)` .

- **Test Accuracy**: 94.35% | **5-Fold CV Accuracy**: 94.75% | **Macro F1**: 0.9459 .

- **Top 3 Deciding Stats**: Clutch Match-Winner Index (12.8%), Batting Impact Index (11.6%), Player of the Match awards (7.8%).
-

3.3 Decision Tree Classifier (CART)

- **WHAT is it?**

A visual flowchart of simple "IF-THEN" rules:

- Rule 1: Is `clutch_match_winner_index` > 24.5?
- If Yes: Is `batting_impact_index` > 62.0? $\rightarrow$ **Elite Tier**
- If No: Is `wickets_taken` < 8? $\rightarrow$ **Developing Tier**

- **WHEN is it used?**

Used when coaches or non-technical executives want a clear, rule-by-rule explanation without math.

- **WHY was it chosen?**

It is the most interpretable model in machine learning. Anyone can trace the decision path on a sheet of paper.

- **HOW does it work & results?**

- Uses the Gini Impurity formula to pick questions that cleanly split the classes. Max depth is set to 6 to prevent memorizing the data.

- **Test Accuracy:** 87.90% | **Macro F1:** 0.8812 .

3.4 Logistic Regression (Multinomial / Softmax)

- **WHAT is it?**

A linear model that outputs probability percentages for each tier (e.g., 85% Elite, 12% Core, 3% Developing).

- **WHEN is it used?**

Used as the baseline probabilistic classifier.

- **WHY was it chosen?**

Gives team owners exact risk percentages during auction bidding wars rather than just a flat label.

- **HOW does it work & results?**

- Uses the Softmax formula to turn raw scores into probabilities that sum to 100%.

- **Test Accuracy:** 87.90% | **5-Fold CV Accuracy:** 92.53% | **Macro F1:** 0.8853 .

4. Unsupervised Learning: Clustering (Tactical Archetypes)

In clustering, the computer has **no labels or answers**. It looks at all 619 players and automatically groups them by similar playing styles.

4.1 K-Means Clustering ($k=5$)

- **WHAT is it?**

An algorithm that groups 619 cricketers into **5 tactical clusters** based on how close their stats are to 5 cluster centers (centroids).

- **WHEN is it used?**

Powers **Tab 4 (Tactical Archetypes & Roster Balance)** in the web app.

- **WHY was it chosen?**

Nominal labels like "Batter" or "Bowler" are too simple for modern T20 cricket:

- Both Virat Kohli and Andre Russell are listed as "Batters", but Kohli anchors the innings while Russell hits death-overs sixes.
- Both Jasprit Bumrah and Yuzvendra Chahal are "Bowlers", but Bumrah bowls yorkers at 145 km/h while Chahal spins middle-over webs.
- K-Means automatically discovered these **5 real tactical archetypes**:

1. **Cluster 0: Tactical Anchor & Top-Order Accumulator** (builds partnerships, high average).
 2. **Cluster 1: High-Impact Pace Spearhead & Death Bowler** (fast yorkers, low death economy, wickets).
 3. **Cluster 2: Explosive Death-Over Finisher & Boundary Hitter** (strike rate > 155, high sixes).
 4. **Cluster 3: Mystery / Control Spin Maestro** (high dot ball %, economical middle overs).
 5. **Cluster 4: Elite Dual-Threat All-Rounder** (contributes heavily with both bat and ball).
- **HOW does it work?**
 1. We tested cluster counts from $k=2$ to $k=7$ using the **Elbow Curve** (inertia) and **Silhouette Score** (cluster separation).
 2. $k=5$ gave the clearest real-world cricket separation ($s = 0.285$).
 3. Saved as `models/kmeans_model.joblib`.
-

5. Dimensionality Reduction (Visualizing 29 Stats in 2D)

5.1 Principal Component Analysis (PCA, 2 Components)

- **WHAT is it?**

A mathematical tool that compresses 29 different statistics down into **2 coordinates (Component 1 on X-axis, Component 2 on Y-axis)** so that all 619 players can be shown on a single 2D scatter plot.
 - **WHEN is it used?**

Powers the interactive 2D Tactical Scatter Plot in **Tab 4** of the web app.
 - **WHY was it chosen?**

Human eyes cannot look at a 29-column table and understand player patterns. PCA keeps the most important information while letting coaches see every player's position on one screen.
 - **HOW does it work & results?**
 - Finds the two directions where player statistics vary the most:
 - **Component 1 (X-axis, 38.16% variance)**: Measures career volume and overall match impact.
 - **Component 2 (Y-axis, 22.63% variance)**: Measures bowling vs. batting specialization (bowlers go up, batters go down, all-rounders stay in the middle).
 - **Together, these 2 coordinates preserve 60.79% of all information** in the original 29 columns.
 - Saved as `models/pca_model.joblib`.
-

6. Model Evaluation & Validation Methods

To make sure our models work reliably on new matches, we use standard academic testing methods:

1. 80/20 Stratified Split:

- 495 players (80%) are used to train the models.
- 124 players (20%) are held out to test the models.
- "Stratified" guarantees the same percentage of Elite, Core, and Developing players in both sets.

2. 5-Fold Stratified Cross-Validation:

- Splits training data into 5 equal parts. The model trains on 4 parts and tests on the 5th part, repeating 5 times.
- We report the average and standard deviation (e.g. 0.9475 ± 0.0142) to prove results are consistent.

3. Core Evaluation Metrics:

- **Accuracy:** Overall percentage of correct tier predictions. KNN achieved **95.16%**.
- **Macro F1-Score:** Harmonic mean of precision and recall, balancing small and large tiers equally.
- **Precision:** Of all players predicted as Elite, how many truly were Elite.
- **Recall:** Of all actual Elite players, how many did the model identify.
- **Confusion Matrix:** A 3 *times* 3 table showing exact correct predictions vs. mistakes.

7. Master Summary Table

Model Name	Task	Syllabus Module	Test Score	5-Fold CV Score	Why It Was Chosen
K-Nearest Neighbors (KNN, k=7)	Talent Tiers	Module V	95.16\% Acc (F1: 0.9501)	<i>0.9030 \pm 0.0221</i>	Accuracy Winner: Classifies players by finding historical peer matches.
Random Forest Classifier	Talent Tiers	Module VIII	94.35\% Acc (F1: 0.9459)	0.9475 \pm 0.0142	Production Winner: Stable across folds and provides feature importance rankings.
Logistic Regression	Talent Tiers	Module V	87.90\% Acc (F1: 0.8853)	<i>0.9253 \pm 0.0180</i>	Linear probabilistic model; gives exact risk percentages.
Decision Tree (CART)	Talent Tiers	Module V	87.90\% Acc (F1: 0.8812)	<i>0.8990 \pm 0.0195</i>	100% human-readable "IF-THEN" flowchart for coaches.
K-Means Clustering	Tactical Styles	Module VII	5 Clusters (<i>s = 0.285</i>)	Validated via Elbow Curve	Discovered 5 real playing styles beyond simplistic Batter/Bowler tags.
PCA	2D Visualization	Module VIII	60.79\% Var (2 Components)	SVD Decomposition	Compresses 29 stats into an interactive 2D scatter plot.